



**UNIVERSITÀ
DI TRENTO**

Data Mining 2025/2026

Corrao Jacopo
Matricola: 258412

Professore: Scettri Giacomo

Indice

1	Introduction	2
2	Data Preparation	3
2.1	Dataset Source	3
2.2	Data Acquisition and engineering	3
2.3	Computational Framework Selection	3
2.4	Schema-on-Read Definition	3
2.5	Data Project and Filtering	4
2.6	Self-Loop Removal	5
2.7	Graph Reconstruction	6
2.8	Optimization	6
3	Methodology and Algorithms	7
3.1	Descriptive Network Analysis: Degree Centrality	7
3.2	Structural Authority: Distributed PageRank	7
3.2.1	Implementation Details	7
3.3	Community Detection: Label Propagation (LPA)	8
3.3.1	Structural Definition of Genres (Unsupervised Learning)	8
4	Results	9
4.1	Network Statistics and Validation	9
4.2	PageRank Authority Rankings	10
4.3	Volume vs. Authority: The Hidden Influencers	10
4.4	Musical Genealogy Networks	11
4.4.1	Top 15 Most-Sampled Artists Network	12
4.5	Artist Classification: Hub Analysis	12
4.6	External Validation Against Ground Truth	14
4.7	Authority Context: Internal vs External Influence	15
4.8	Graph Statistics Summary	16
4.8.1	Interactive D3.js Visualization	17
4.9	Community Detection and Genealogical Families	17
4.9.1	Macro-Flow: Inter-Community Genealogy	18
5	Discussion and Interpretation	20
5.1	Theoretical Implications	20
5.2	Methodological Contributions	20
5.3	Limitations and Future Work	20
6	Conclusion	22

1 Introduction

The music industry typically measures success through sales charts and streaming numbers (popularity). However, cultural impact is structurally different from popularity. A "one-hit wonder" might sell millions, but an obscure funk track from the 1970s might be sampled by hundreds of hip-hop artists, forming the backbone of an entire genre. The objective of this project is to shift the analysis from volume to structure. By modeling the music ecosystem as a **Directed Graph** where **nodes** are **songs** and **edges** represent **sampling relationships** we aim to:

- Identify the structural "ancestors" of modern music (Authority).
- Detect communities of influence that transcend traditional genre labels.
- Analyze the flow of influence using distributed graph algorithms.

2 Data Preparation

2.1 Dataset Source

MusicBrainz Data was sourced from the **MusicBrainz** Database, an open music encyclopedia that collects user-contributed metadata. Unlike simplified datasets found on platforms like Kaggle, **MusicBrainz** provides raw PostgreSQL database dumps. The complexity of this dataset lies in its highly normalized structure (Third Normal Form), which requires significant data engineering to reconstruct meaningful relationships. We utilized the core tables: `recording` (songs), `artist_credit` (artist names), `l_recording_recording` (relationships), and the bridging table `link`, which connects specific relationships to their semantic definitions found in the `link_type` table.

2.2 Data Acquisition and engineering

The raw data consisted of tab-separated value (TSV) files without headers, extracted from the `mbdump.tar.bz2` archive. The transformation process from raw relational tables to a structured property graph involved three key phases: Schema Definition, Projection, and Graph Reconstruction.

2.3 Computational Framework Selection

Before addressing the data schema, it is necessary to justify the architectural choice. Given the relational complexity of the MusicBrainz database (highly normalized) and the potential size of the graph, traditional single-node tools like Pandas would be inefficient for the required multi-way joins and iterative computations.

Apache Spark (PySpark) was chosen as the framework to leverage:

- **Distributed Processing:** To handle data ingestion and transformation in parallel.
- **Lazy Evaluation:** To optimize the execution plan of complex transformation pipelines.
- **Iterative Computation:** Essential for implementing graph algorithms like PageRank from scratch.

2.4 Schema-on-Read Definition

Having established the framework, the first challenge was ingestion. Since the raw files lacked headers, relying on automatic schema inference was computationally

expensive and error-prone (risking incorrect data type assignment). I manually defined the StructType for each table based on the MusicBrainz documentation. This ensures strict typing and robust data ingestion.

```
recording_schema = StructType([
    StructField("id", IntegerType(), True),
    StructField("gid", StringType(), True),
    StructField("name", StringType(), True),
    StructField("artist_credit", IntegerType(), True),
    StructField("length", IntegerType(), True),
    StructField("comment", StringType(), True),
    StructField("edits_pending", IntegerType(), True),
    StructField("last_updated", StringType(), True),
    StructField("video", StringType(), True)
])
```

Listing 1: Code Snippet: Manual Schema Definition

2.5 Data Project and Filtering

To optimize memory usage on the Spark driver and executors, I applied column pruning immediately after loading. Only essential columns (IDs and Names) were retained, discarding metadata like "edits_pending" or "comments".

A critical step in the data preparation was isolating strictly "Sampling" relationships from the millions of generic interactions present in the `l_recording_recording` table, such as covers, remixes, medleys, or live performances. Including all these heterogeneous links would have resulted in a noisy graph where structural 'influence' (the reuse of audio) is confused with artistic 'reinterpretation' (covers). To achieve a precise genealogy—defined as instances where a piece of audio is physically reused—I performed an exploratory query on the `link_type` metadata table to identify the specific Primary Keys associated with sampling descriptions.

This exploration identified two distinct IDs:

- **ID 69:** "Samples material" (Legacy type)
- **ID 231:** "Is sampled by" / "Samples material" (Current type)

Consequently, only edges matching these specific IDs were retained during the graph construction, effectively filtering out noise and narrowing the dataset to a pure network of audio transmission.

2.6 Self-Loop Removal

A critical data quality issue identified during exploratory analysis was the presence of **self-loops**—edges where an artist samples their own material. These relationships totaled 1,407 instances (approximately 6.4% of all edges) and typically represent:

- **Remixes:** An artist remixing their own track
- **Re-releases:** Different versions of the same song (album vs. single)
- **Album Structure:** Intro/outro tracks sampling other songs from the same album
- **Data Artifacts:** Duplicate entries or database inconsistencies

Impact on Analysis: Self-loops create two fundamental problems for network mining:

1. **Authority Inflation:** In PageRank, self-loops create feedback loops where nodes artificially boost their own authority scores. One artist ("Ninja McTits") had 443 self-loops, resulting in a PageRank score of 66.45—dramatically higher than the legitimate top artist (James Brown: 45.44).
2. **Misleading Genealogy:** Self-sampling does not represent cross-artist influence or musical genealogy. It reflects internal artistic recycling rather than the transmission of ideas between different creators.

Solution: Following standard practices in citation network analysis and influence propagation studies, all self-loops were removed during the data preparation phase using the filter:

```
df_graph = df_graph.filter(  
    col("Sampler_Artist_Name") != col("Original_Artist_Name")  
)
```

Listing 2: Code Snippet: Self-Loop Removal

This filtering reduced the edge count from 22,135 to 20,728 edges at the artist level (after collapsing multiple song-level edges between the same artist pair). The final graph at the song level contains 58,621 sampling events—each representing a specific recording-to-recording relationship—which are then aggregated into weighted artist-level edges for the PageRank and clustering analyses. This

distinction is important: the 58,621 figure reflects fine-grained song-level interactions, while the 20,728 figure represents the collapsed artist-to-artist topology used for structural analysis. Both levels of granularity are preserved in the processed dataset.

2.7 Graph Reconstruction

The MusicBrainz database is highly normalized. A "Sampling" relationship is stored as a numeric link between two recording IDs (`entity0` and `entity1`), without any text information about the song title or artist name. To build a human-readable graph (e.g., Artist A samples Artist B), I implemented a de-normalization pipeline involving a chain of joins across four DataFrames:

- **Linkage:** Join `l_recording_recording` with the `link` table. This step was necessary to access the `link_type` attribute and filter strictly for the IDs identified in the exploration phase (69, 231).
- **Source Enrichment:** Join the result with `recording` and `artist_credit` to resolve the Source (Child) song and artist names.
- **Target Enrichment:** Repeat the join with `recording` and `artist_credit` to resolve the Target (Parent) song and artist names.

2.8 Optimization

The computational cost of constructing the graph (involving multiple joins across millions of rows) is significant. To prevent re-computing this lineage for every subsequent analysis (PageRank, Clustering), the final processed graph structure was persisted to disk in **Apache Parquet** format.

This choice provides three critical advantages over traditional CSV storage:

- **Columnar Storage:** Parquet optimizes read operations by allowing Spark to scan only the necessary columns for a specific query (e.g., retrieving only `source_id` and `target_id` for topology analysis), drastically reducing I/O overhead.
- **Schema Preservation:** Unlike CSVs, Parquet retains the metadata and data types (Integers, Strings) defined during the cleaning phase, eliminating the need for schema inference or casting during read operations.
- **Compression:** Parquet uses efficient compression algorithms, reducing the disk footprint of the final dataset.

3 Methodology and Algorithms

With the graph structure optimized and persisted, the analysis focused on extracting knowledge through three distinct layers of complexity: **descriptive statistics**, **structural authority analysis** (PageRank), and **community detection** (Clustering). Crucially, to fully leverage the distributed nature of Spark and demonstrate a deep understanding of graph theory, I implemented the iterative graph algorithms from scratch using PySpark transformations (MapReduce paradigm) rather than relying on pre-built libraries like GraphFrames.

3.1 Descriptive Network Analysis: Degree Centrality

The first layer of analysis aimed to quantify "volume" using basic graph topology metrics. Using Spark SQL aggregations, I calculated the Degree Centrality for each node:

- **In-Degree (Most Sampled)**: Represents the number of incoming edges. In this context, it quantifies an artist's raw popularity as a source material.
 - Implementation:

```
groupBy("Original_Artist_Name").count()
```

- **Out-Degree (Top Samplers)**: Represents the number of outgoing edges. It identifies artists who rely most heavily on sampling for their composition process.

3.2 Structural Authority: Distributed PageRank

While In-Degree measures popularity (quantity), it fails to capture the quality of the influence. To solve this, I modeled the "**Authority**" of an artist using the **PageRank algorithm**. Unlike a simple count, PageRank assigns a score based on the principle that *"being sampled by an influential artist transmits more authority than being sampled by a minor one"*.

3.2.1 Implementation Details

I implemented the algorithm manually via an iterative MapReduce process with **convergence criterion**.

The logic follows the standard formula:

$$PR(A) = (1 - d) + d \sum \frac{PR(B)}{OutDegree(B)}$$

where d is the damping factor (set to 0.85).

The implementation ensures that authority flows correctly: edges point from Sampler to Original Artist, so that the Original Artist (being sampled) receives authority through *incoming* edges. The algorithm converged in **11 iterations** with tolerance < 0.0001 , demonstrating efficient convergence.

The PySpark implementation addressed specific technical challenges:

- **Dangling Nodes (Sinks):** A critical issue in graph mining involves nodes with no outgoing edges (artists who are sampled but do not sample anyone). These nodes act as "black holes" for the rank score. To prevent rank leakage, I utilized a `right_outer_join` to identify these nodes and manage their score redistribution.
- **Lineage Truncation:** Spark's lazy evaluation builds a growing lineage graph with each iteration, which can lead to `StackOverflowError`. I mitigated this by applying `.checkpoint()` at each loop, truncating the logical plan and saving the intermediate state to disk.

3.3 Community Detection: Label Propagation (LPA)

To identify distinct "musical families", I implemented a **Label Propagation Algorithm (LPA)**. This is an unsupervised clustering technique where nodes adopt the "label" of their neighbors based on majority voting.

Algorithm Logic:

- **Initialization:** Every artist starts with a unique label equal to their own ID.
- **Propagation (6 Iterations):** In each iteration, a "Sampler" node adopts the label of its "Sampled" node. If a node samples multiple artists, the `greatest()` function deterministically selects the dominant label.
- **Convergence:** Over multiple iterations, connected components (chains of sampling) converge to a single shared label—the ID of the "Ancestor" or "Root" artist.

3.3.1 Structural Definition of Genres (Unsupervised Learning)

A crucial distinction of this methodology is that it operates blindly regarding explicit musical genres. We did not import metadata tables containing tags like "Hip Hop" or "Rock". Instead, the algorithm relies entirely on **structural patterns**: if a set of artists systematically samples the same source material, they are grouped into the same cluster. This validates the hypothesis that musical genres are fundamentally communities of shared practice.

4 Results

4.1 Network Statistics and Validation

The final processed graph, after extensive data engineering using MusicBrainz GIDs and removing self-loops (artists sampling themselves), comprises **58,621 sampling events** (edges) connecting **68,527 songs** (nodes) across **23,503 unique artists**.

The network exhibits classic **scale-free** characteristics, as evidenced by the power-law degree distribution shown in Figure 1.




Figure 1: Power-law validation: cumulative degree distribution (log-log scale). The approximately linear trend in log-log space is consistent with a power-law distribution, a hallmark of scale-free networks. Concentration analysis shows top 1% of artists control 23.6%, top 5% control 46.6%, and top 10% control 58.2% of all sampling events.

Key Network Properties:

- **Graph Density:** 0.00001248 (sparse network, as expected in real-world influence networks)

- **Mean In-Degree:** 3.50 (average times an artist is sampled)
- **Mean Out-Degree:** 6.02 (average samples used per artist)
- **Maximum In-Degree:** 371 (Daft Punk - most sampled artist)
- **Power-Law Concentration:** Top 1% of artists control 23.6% of all sampling events, Top 10% control 58.2%

This concentration pattern validates the scale-free network hypothesis: musical influence follows preferential attachment, where established authorities accumulate disproportionate influence over time.

4.2 PageRank Authority Rankings

The PageRank algorithm converged in **11 iterations** (tolerance < 0.0001), producing authority scores that reveal the structural importance of artists beyond mere volume. Table 1 presents the comparison of the top 10 artists by volume versus authority.

Tabella 1: Comparison of Top 10 Artists by Volume vs. Authority

By Volume			By Authority		
Rank	Artist	Count	Rank	Artist	Score
1	Daft Punk	371	1	Daniel Ingram	59.3799
2	Lady Gaga	282	2	James Brown	45.4355
3	The Beatles	268	3	2Pac	44.8128
4	JAY-Z	260	4	音部	38.5237
5	PSY	253	5	外神田文芸高校	33.7655
6	Linkin Park	230	6	Daft Punk	27.1559
7	Michael Jackson	222	7	Lyn Collins	27.0612
8	Beastie Boys	209	8	室Pixel	26.0531
9	James Brown	204	9	Kraftwerk	25.8983
10	Katy Perry	187	10	John Lennon	25.2695

4.3 Volume vs. Authority: The Hidden Influencers

Comparing the results of Degree Centrality (Volume) against PageRank (Authority) reveals significant insights into the music ecosystem, as shown in Figure 2.




Figure 2: Top 20 artists: Volume (In-Degree) vs. Authority (PageRank). The two rankings reveal a critical gap: high volume does not always translate to high structural authority, exposing the role of community density in PageRank scores.

- **The Volume Kings:** Modern electronic and pop icons like **Daft Punk** (371 events) and **Lady Gaga** (282) dominate the volume metrics. These are the artists most frequently sampled in absolute terms.
- **The Classic Authorities:** **James Brown**, **Lyn Collins**, and **Kraftwerk** rank high in PageRank authority despite lower raw volume, indicating that their samplers are themselves highly influential artists — a hallmark of genuine structural importance.
- **The Community-Driven Outlier:** **Daniel Ingram** (composer for *My Little Pony*) achieves the highest PageRank (59.38) despite only 130 sampling events. As explored later in the Authority Context analysis, this is driven by a dense, insular fan community rather than mainstream influence — illustrating how PageRank can be amplified by tight cluster structures.

4.4 Musical Genealogy Networks

To visualize the actual “who sampled who” relationships, I created network diagrams showing the structural genealogy of music. Each visualization isolates a specific aspect of the network to make the results clear and interpretable.

4.4.1 Top 15 Most-Sampled Artists Network

Figure 3 shows the core sampling relationships among the 15 most-sampled artists by volume (in-degree). To avoid visual clutter, only connections with weight ≥ 2 are displayed.




Figura 3: Top 15 most-sampled artists and their inter-connections. Arrows point from sampler to original artist. Node size is proportional to sampling volume. This view reveals which top artists also sample each other, forming the backbone of the sampling ecosystem.

4.5 Artist Classification: Hub Analysis

Figure 4 classifies artists based on their sampling behavior, plotting in-degree (times sampled) against out-degree (samples used). The dashed lines represent median values, dividing the space into four quadrants that reveal distinct artist roles.




Figure 4: Hub analysis scatter plot. X-axis: out-degree (samples used). Y-axis: in-degree (times sampled). Top 15 most-sampled artists are highlighted in red. Four quadrants are delineated by median lines: Authorities (top-left), Bridges (top-right), Heavy Samplers (bottom-right), and Peripheral Artists (bottom-left).

Artist Classifications:

- **Pure Authorities** (top-left quadrant): Artists like James Brown and Daft Punk — sampled extensively but use few samples themselves. They are the “source material” of modern music.
- **Bridges** (top-right quadrant): Artists who both sample heavily AND get sampled. As discovered in our bridge analysis (Figure 5), evolutionary nodes like **JAY-Z**, **Beastie Boys**, and **Ye (Kanye West)** maintain a near-perfect balance, acting as the crucial connective tissue of modern music.
- **Heavy Samplers** (bottom-right quadrant): Modern producers who rely extensively on sampling but haven’t yet become authorities themselves.

- **Peripheral Artists** (bottom-left quadrant): Artists with below-median in-degree and out-degree. These represent the majority of nodes in the long tail of the distribution — niche creators or artists with few documented sampling relationships in the dataset.




Figura 5: Top 15 Evolutionary Bridges. These artists maintain the highest geometric mean of in-degree and out-degree, representing key evolutionary nodes in the musical network.

4.6 External Validation Against Ground Truth

To assess the robustness of our PageRank authority rankings, we performed external validation using a curated list of the 10 most-sampled artists from **WhoSampled**, an independent sampling reference. We computed the Spearman rank correlation between the WhoSampled ground-truth ordering and our PageRank rankings, considering only the same 10-artist subset.

The resulting Spearman correlation coefficient is $\rho = 0.37$, indicating a **moderate positive correlation**. While not perfect, this is notable given fundamental methodological differences: WhoSampled ranks by raw sampling count, whereas PageRank measures structural influence. The moderate agreement confirms that our authority metric captures genuine influence while diverging in meaningful

ways—for example, PageRank elevates structurally influential artists like Lyn Collins (ground-truth rank 5, system rank 2) and Kraftwerk (ground-truth rank 6, system rank 3) above high-volume but less structurally central artists like Madonna (ground-truth rank 10, system rank 9) and Public Enemy (ground-truth rank 2, system rank 10).

Full match results are saved in `outputs/external_validation.csv`. This validation provides independent confirmation that our PageRank implementation identifies genuinely influential artists while offering a dimension of authority distinct from raw count-based metrics.

4.7 Authority Context: Internal vs External Influence

While PageRank effectively models authority, it does not distinguish between *global* mainstream influence and *insular* community dominance. To investigate the "Daniel Ingram anomaly" (an artist with moderate volume but the highest PageRank), we analyzed the composition of incoming samples for the top 15 authorities.

Figure 6 proves the "Insular Community Hypothesis". While mainstream authorities like James Brown receive the vast majority of their influence from outside their immediate community, Daniel Ingram's authority is driven by a dense, internal network of dedicated remixers.




Figura 6: Authority Context: Internal vs External Influence. This chart reveals whether an artist’s authority is driven by a massive, isolated sub-culture (Internal) or by universal mainstream adoption across genres (External).

4.8 Graph Statistics Summary

Table 2 provides a comprehensive overview of the network’s macroscopic properties.

Metric	Value
Graph Structure	
Total Sampling Events (Edges)	58,621
Unique Songs (Nodes)	68,527
Unique Artists	23,503
– Artists Who Sample	9,709
– Artists Being Sampled	16,634
In-Degree Statistics	
Mean	3.50
Std Dev	10.56
Maximum	371
Out-Degree Statistics	
Mean	6.02
Std Dev	33.20
Maximum	1,488

Tabella 2: Summary of key graph statistics. The transition to robust GID-based entity resolution resulted in a vastly expanded network connecting over 68,000 songs.

4.8.1 Interactive D3.js Visualization

To effectively explore the immense scale of the expanded dataset (over 58,000 edges), I also developed an interactive web-based D3.js visualization that allows users to pan, zoom, and click on nodes to inspect real-time Authority Scores and isolate specific sampling lineages. The interactive visualization is available at: interactive-genealogy.html

4.9 Community Detection and Genealogical Families

The Label Propagation Algorithm detected **13,855 distinct clusters** with an intra-cluster edge fraction (simplified modularity) of **0.1926**, indicating a complex and overlapping community structure. This metric represents the proportion of edges that fall within the same cluster; a value of 0.19 means that approximately 19% of sampling connections occur between artists in the same genealogical community, while 81% cross community boundaries. This suggests the music network is highly interconnected across genre boundaries.




Figure 7: Distribution of cluster sizes. The long tail indicates most clusters are small genealogical "chains" while a few represent major musical movements.

Cluster Analysis:

- **Largest Cluster:** Contains 675 songs, representing the dominant, highly integrated core of modern sampling.
- **Simplified Modularity 0.1926:** 19% of edges are intra-cluster, indicating a low but detectable community structure. This is expected in modern digital music where producers freely mix samples from disparate sources.
- **13,855 clusters:** The vast number of clusters highlights the fragmented nature of niche internet subcultures and isolated sampling chains that haven't connected to the main network.

4.9.1 Macro-Flow: Inter-Community Genealogy

While the Label Propagation Algorithm identified isolated communities, music evolution occurs when these communities interact. By aggregating edges at the cluster level, we mapped the macroscopic flow of influence.

As shown in Figure 8, distinct musical pipelines emerge. Instead of analyzing individual artists, we can observe how entire ecosystems feed into others, providing a higher-level view of genre cross-pollination.




Figura 8: Macroscopic Sampling Flow Between Communities. Node size represents total incoming cluster influence, and edge thickness represents the volume of sampling between distinct communities.

5 Discussion and Interpretation

5.1 Theoretical Implications

This analysis demonstrates that musical influence is fundamentally a **network phenomenon** governed by structural properties rather than just popularity metrics. Three key insights emerge:

1. **Authority vs. Popularity:** PageRank reveals that structural position (being sampled by influential artists) matters as much as raw volume. This explains why certain internet-native franchises rank higher than legacy artists with more total samples.
2. **Cross-Genre Genealogy:** The emergence of media franchises (animation, video games) as top authorities demonstrates that sampling creates unexpected genealogical paths beyond traditional genre boundaries.
3. **Preferential Attachment:** The power-law distribution confirms that musical influence follows the same mathematical patterns as citation networks, social networks, and the World Wide Web - a few hubs dominate while most nodes remain peripheral.

5.2 Methodological Contributions

- **From-Scratch Implementation:** Implementing PageRank and Label Propagation manually in PySpark demonstrates deep understanding of distributed graph algorithms beyond using pre-built libraries.
- **Data Engineering:** Successfully navigating MusicBrainz's complex relational schema (Third Normal Form) and transforming it into a graph structure showcases real-world data engineering skills.
- **Convergence Optimization:** Implementing convergence criteria with adaptive stopping improved algorithm efficiency, achieving convergence in just 11 iterations with tolerance < 0.0001 .
- **Visualization Strategy:** Creating clean, non-overlapping network visualizations makes complex structures interpretable.

5.3 Limitations and Future Work

- **Data Quality:** While self-loop removal addressed the most critical data quality issue (artists sampling themselves), the MusicBrainz database still con-

tains inconsistencies such as placeholder artists and duplicate entries. Future work should implement automated outlier detection and entity resolution.

- **Temporal Analysis:** The current analysis is static. Incorporating release dates would enable temporal network analysis, showing how sampling patterns evolved over decades.
- **Genre Metadata:** While Label Propagation is unsupervised, incorporating explicit genre tags would allow validation of whether structural clusters align with conventional genre definitions.
- **Multi-Hop Genealogy:** Extending the analysis to trace multi-generational sampling chains (Original $\rightarrow$ Sample 1 $\rightarrow$ Re-sample) would reveal deeper genealogical paths.

6 Conclusion

This project successfully applied distributed graph mining techniques to uncover the hidden genealogy of modern music. By modeling sampling relationships as a directed graph and implementing PageRank and Label Propagation from scratch in Apache Spark, we revealed structural patterns invisible to traditional volume-based metrics.

Key Findings:

- **Daniel Ingram** and Japanese electronic projects like **Denonbu** emerge as the new unexpected authorities, showcasing the immense structural influence of modern media franchises and internet culture on contemporary sampling.
- **James Brown** and classic funk icons remain foundational pillars, but now share the top echelon with digital-era creators.
- The network follows a strong **power-law distribution**: the top 1% of artists control 23.6% of all sampling events.
- **13,855 genealogical families** detected with an intra-cluster edge fraction of 0.1926, indicating a highly interconnected and overlapping modern music landscape where 81% of sampling crosses community boundaries.
- **Entity Resolution**: Transitioning to MusicBrainz GIDs revealed a vastly larger network (58k+ edges) than previously estimated using brittle string matching.
- **Convergence in 11 iterations** validates the optimized distributed PageRank implementation with proper sink-mass redistribution.

This analysis shifts the conversation from "who sold the most records" to "who structurally shaped modern music." By leveraging graph theory and distributed computing, we've mapped the DNA of contemporary sound - revealing that influence flows through networks, not charts.

Impact: This methodology could extend to other domains: academic citation networks (which papers are foundational?), social media influence (who drives conversation?), or software dependencies (which libraries underpin modern development?). The tools are universal; the insights are profound.